



# Audio data analysis

***CES Data Science***

## Slim ESSID

Audio, Acoustics & Waves Group - Image and Signal Processing dpt.

[slim.essid@telecom-paristech.fr](mailto:slim.essid@telecom-paristech.fr)

<http://www.telecom-paristech.fr/~essid>

## Credits



O. GILLET, C. JODER, N. MOREAU, G. RICHARD, F. VALLET, ...

Slim Essid

# About “audio”...

## ► Audio frequency:

the range of audible frequencies (20 to 20,000 Hz)




[CC Attribution 2.5 Generic](#)

# About “audio”...





## ► Audio content categories

Speech



Music

Environmental




# About “audio”...

## ► An important distinction: speech vs non-speech

### Speech signals

“Simple” production model:  
the **source-filter** model



### Music & non-speech (environmental)

No generic production model:  
“timbre”, “pitch”, “loudness”, ...



*Image: Edward Flemming, course materials for 24.910 Topics in Linguistic Theory: Laboratory Phonology, Spring 2007. MIT OpenCourseWare (<http://ocw.mit.edu/>), Massachusetts Institute of Technology. Downloaded on 05 May 2012*

# About “audio”...



## ► Different research communities

### Music Information Research

Music  
classification  
(genre, mood, ...)

Transcription

Rhythm  
analysis

...

Signal  
representations

Audio coding

Source  
separation

Sound  
synthesis

...

### Speech

Speech  
recognition

Speaker  
recognition

Speech  
enhancement

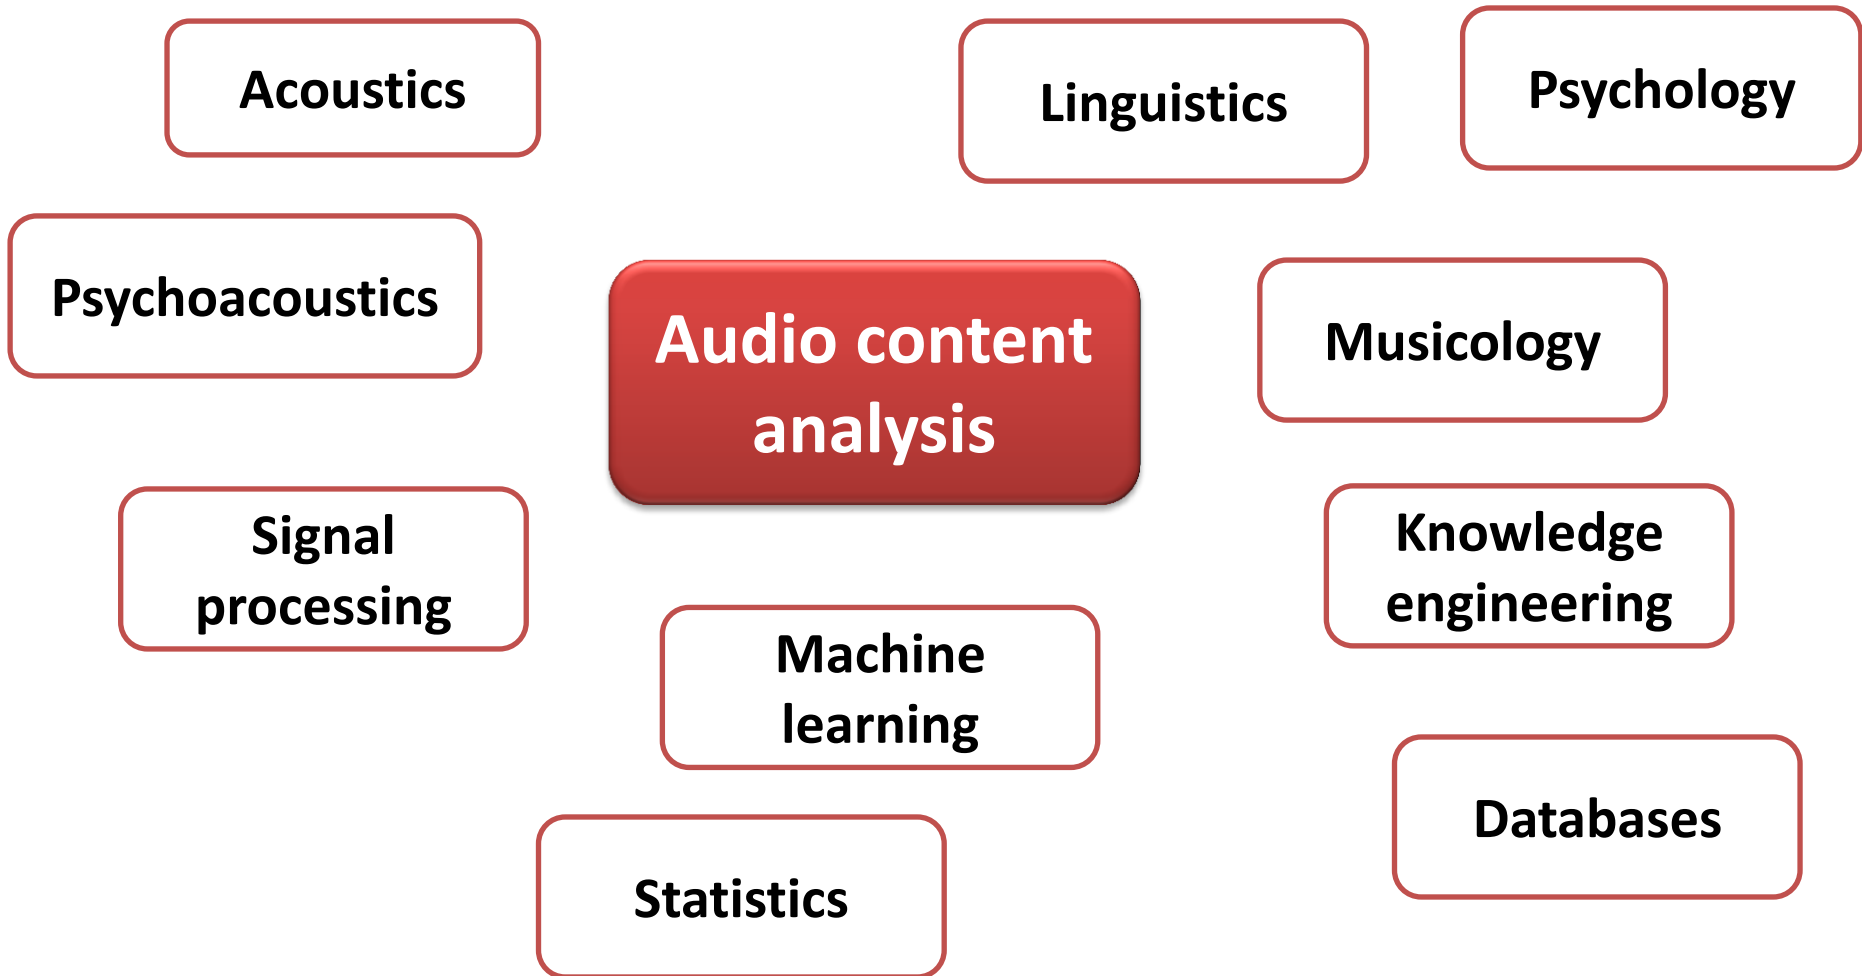
...

## Computer audition

# About “audio”...



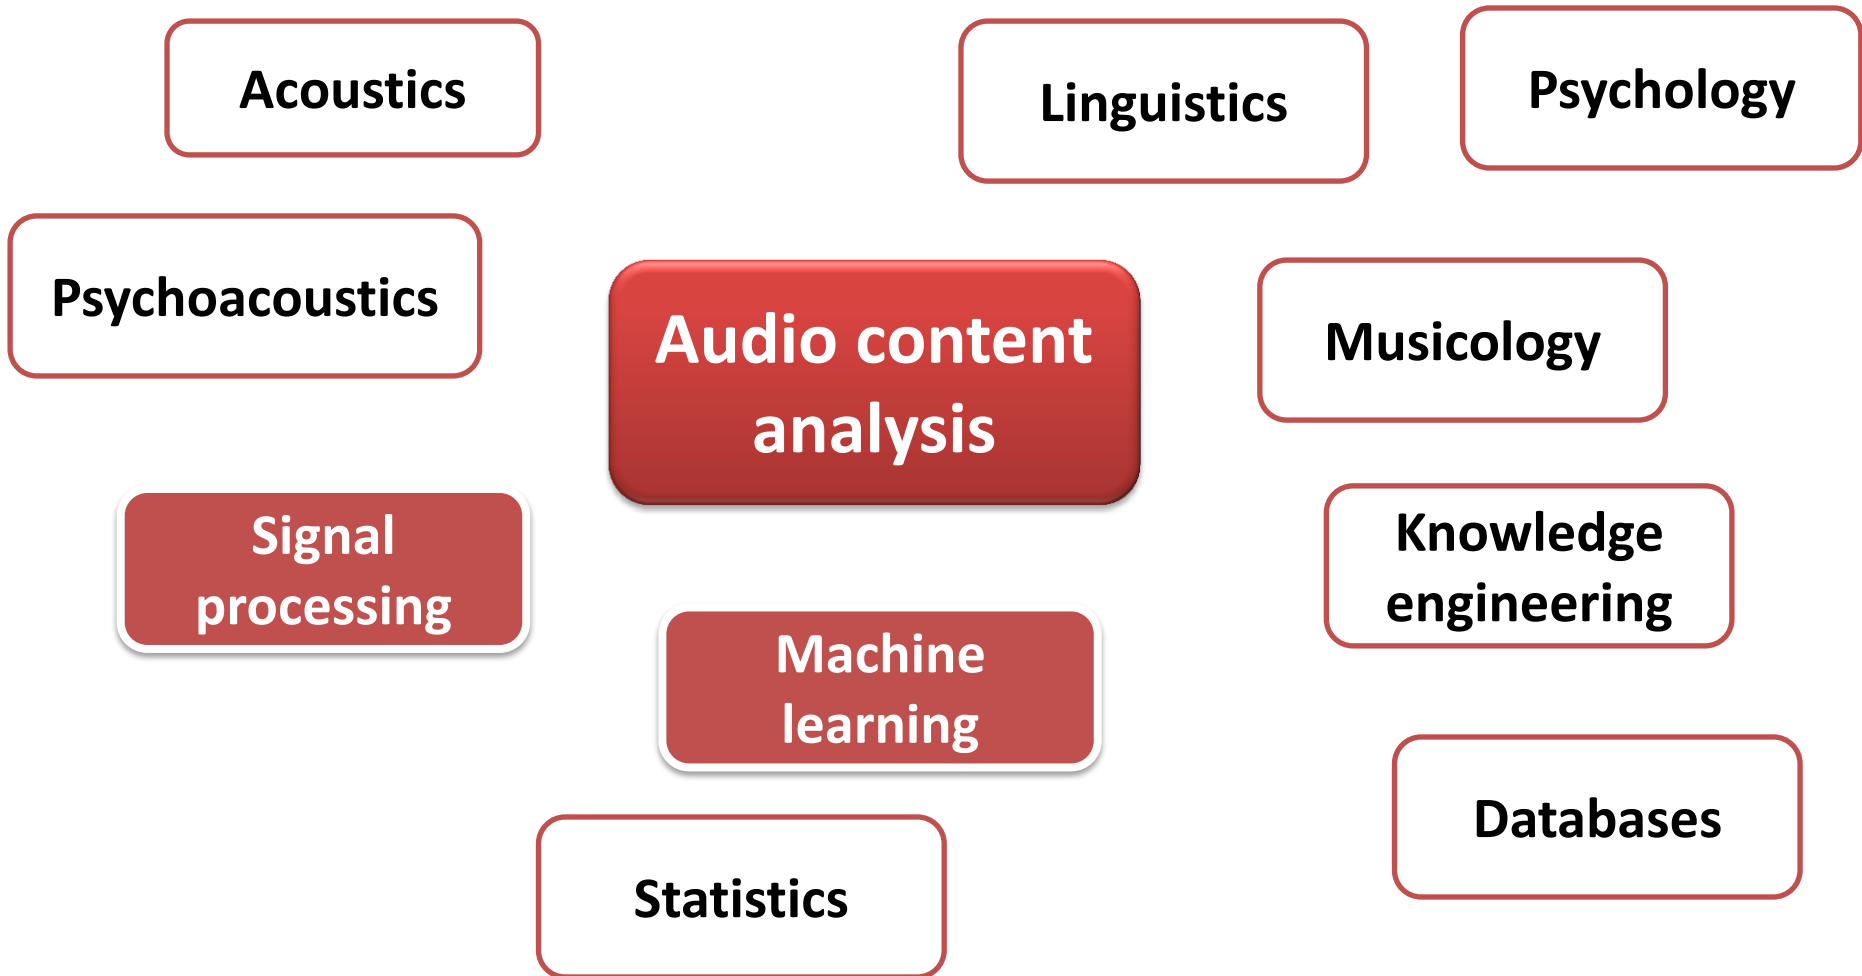
## ► Research fields



# About “audio”...



## ► Research fields







- **For archive management, indexing**

- » Content segmentation and classification: speech/music/jingles..., speakers
- » Music **autotagging**: genre (classical, jazz, rock,...), mood, usage...
- » Search engines

- **For broadcasters**

- » Music/effects/speech excerpt search
- » Playlist generation, Djing





- **For designers and producers**

- » Audio sample search
- » Music transcription (beat, rhythm, chords, notes)
- » Broadcast content monitoring, plagiarism detection, *hit prediction*

- **For end-users**

- » Content-based search (shazam++)
- » Non-linear and interactive content consuming (“skip intro”, “replay the chorus”, Karaoke: “remove the vocals”...)
- » Recommendation
- » Personalised playlist generation




- **Motivation** for audio-driven content analysis
  - » critical information is conveyed by the audio content
  - » audio and visual information play complementary roles for the detection of key concepts/events
- Video examples

# Audio-driven multimedia analysis



## ► Video examples



→ Use audio-based **laughter** detection



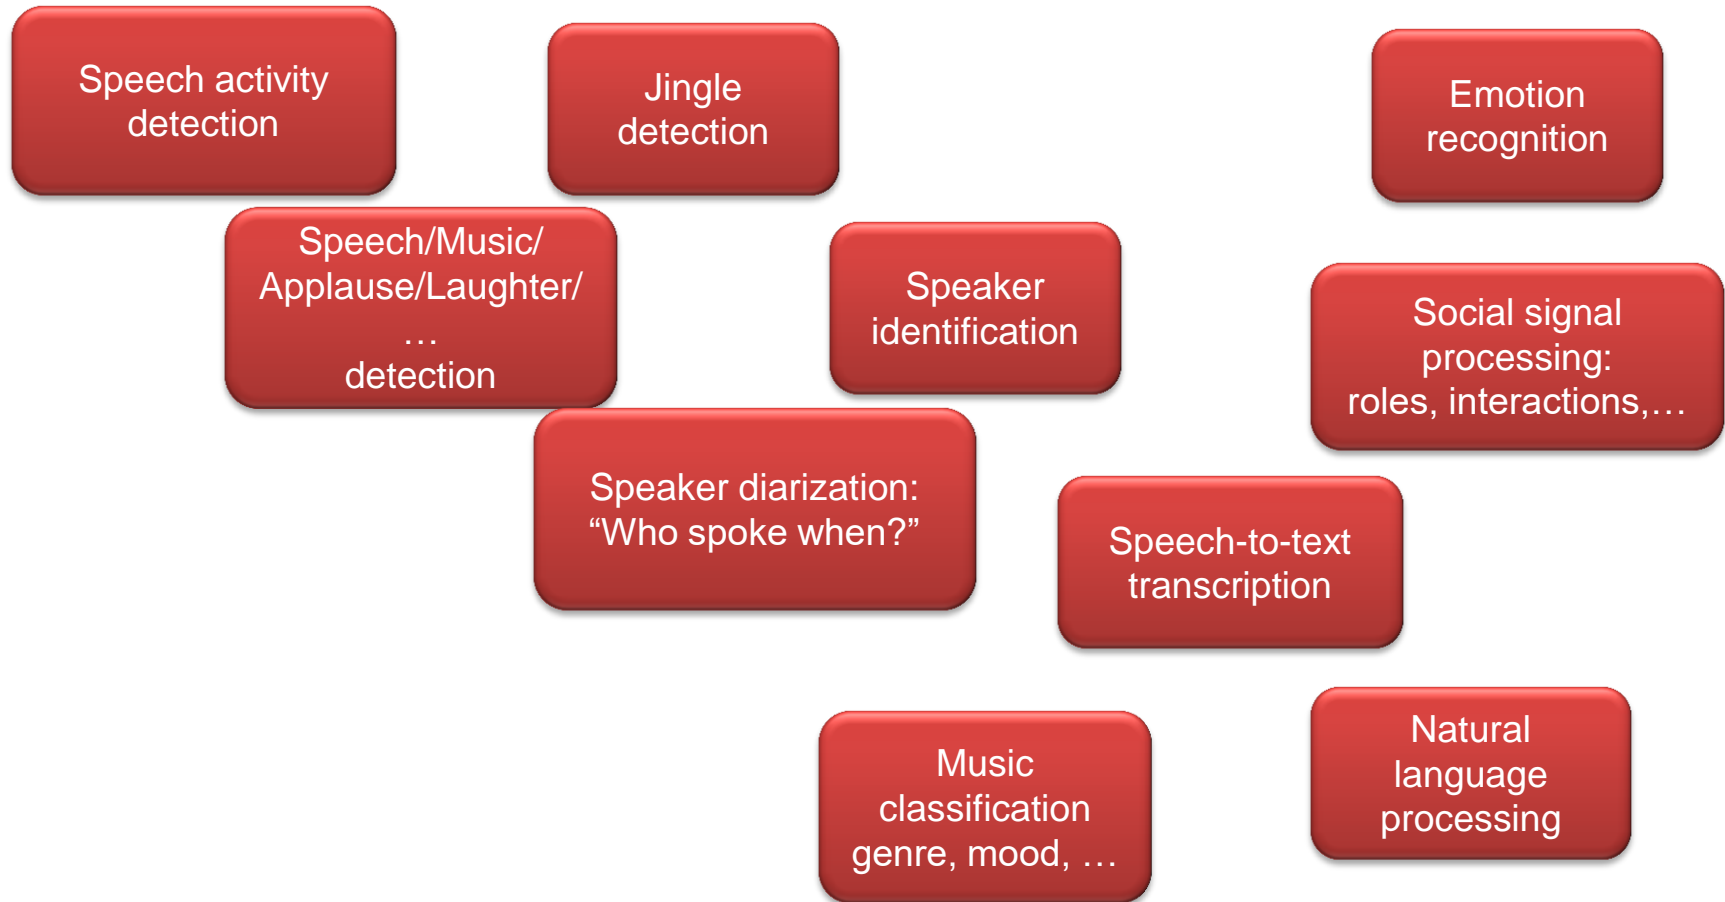
- **Applause** detection
- **Cheering** detection

- Keyword spotting: **“Goal!”**
- Sound **loudness**
- Applause/cheering detection

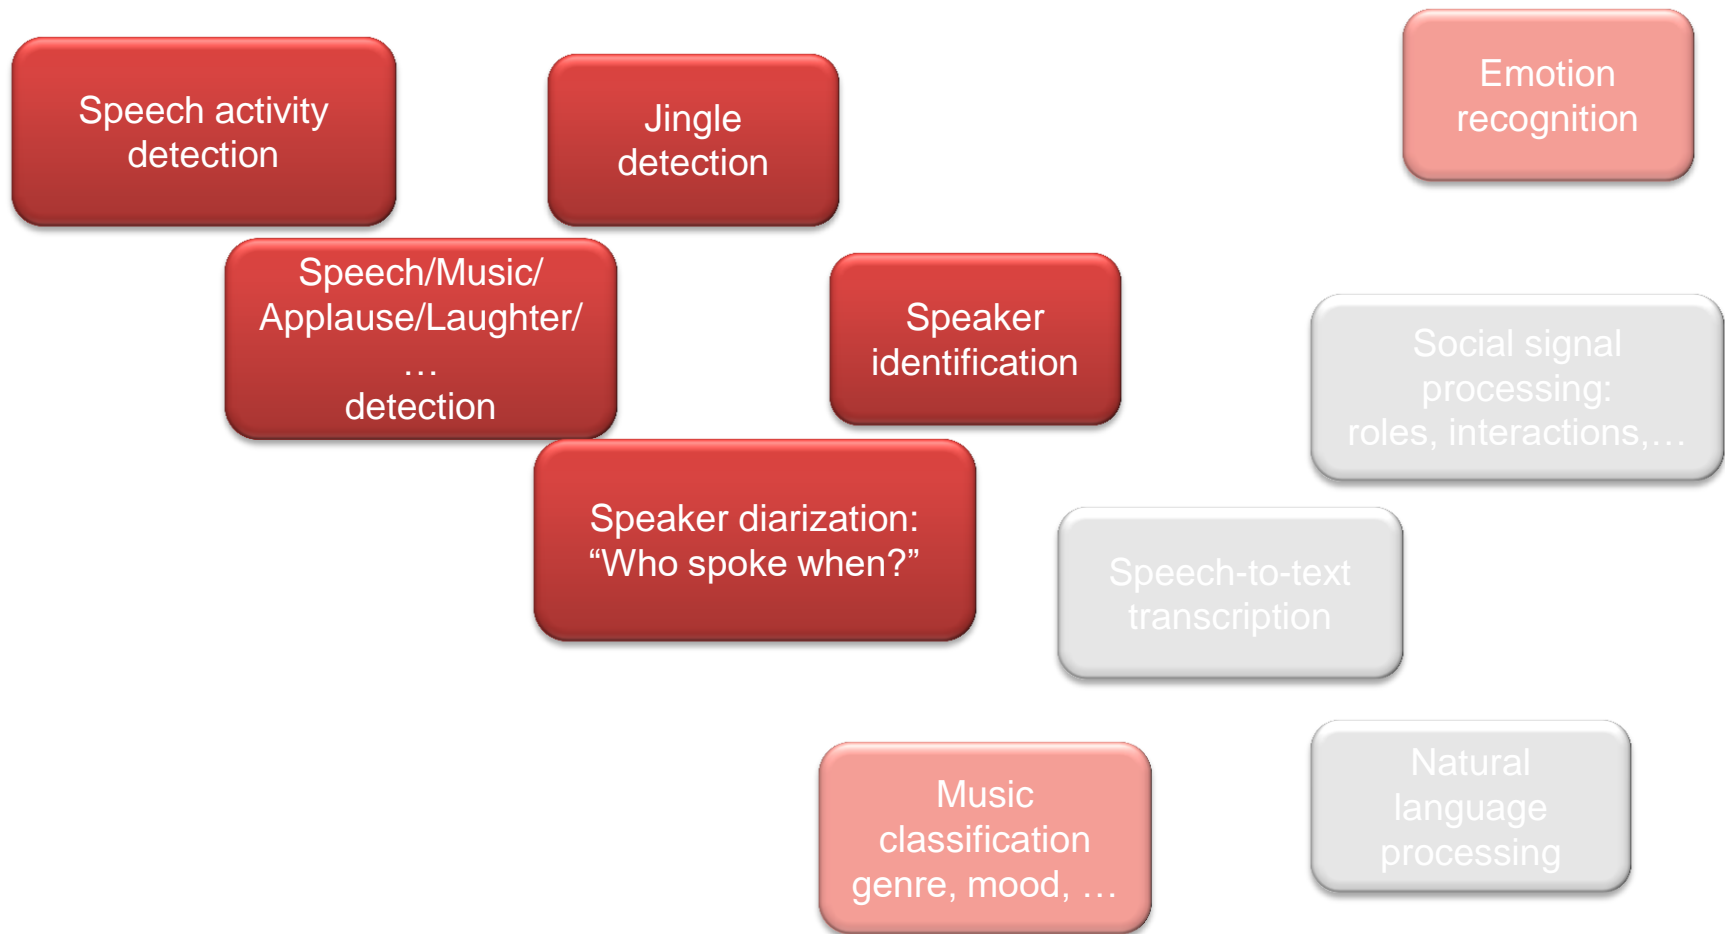
# Audio-driven multimedia analysis



## ► Key audio-based components



→ At the heart of all components: a **classification** task (supervised or unsupervised)

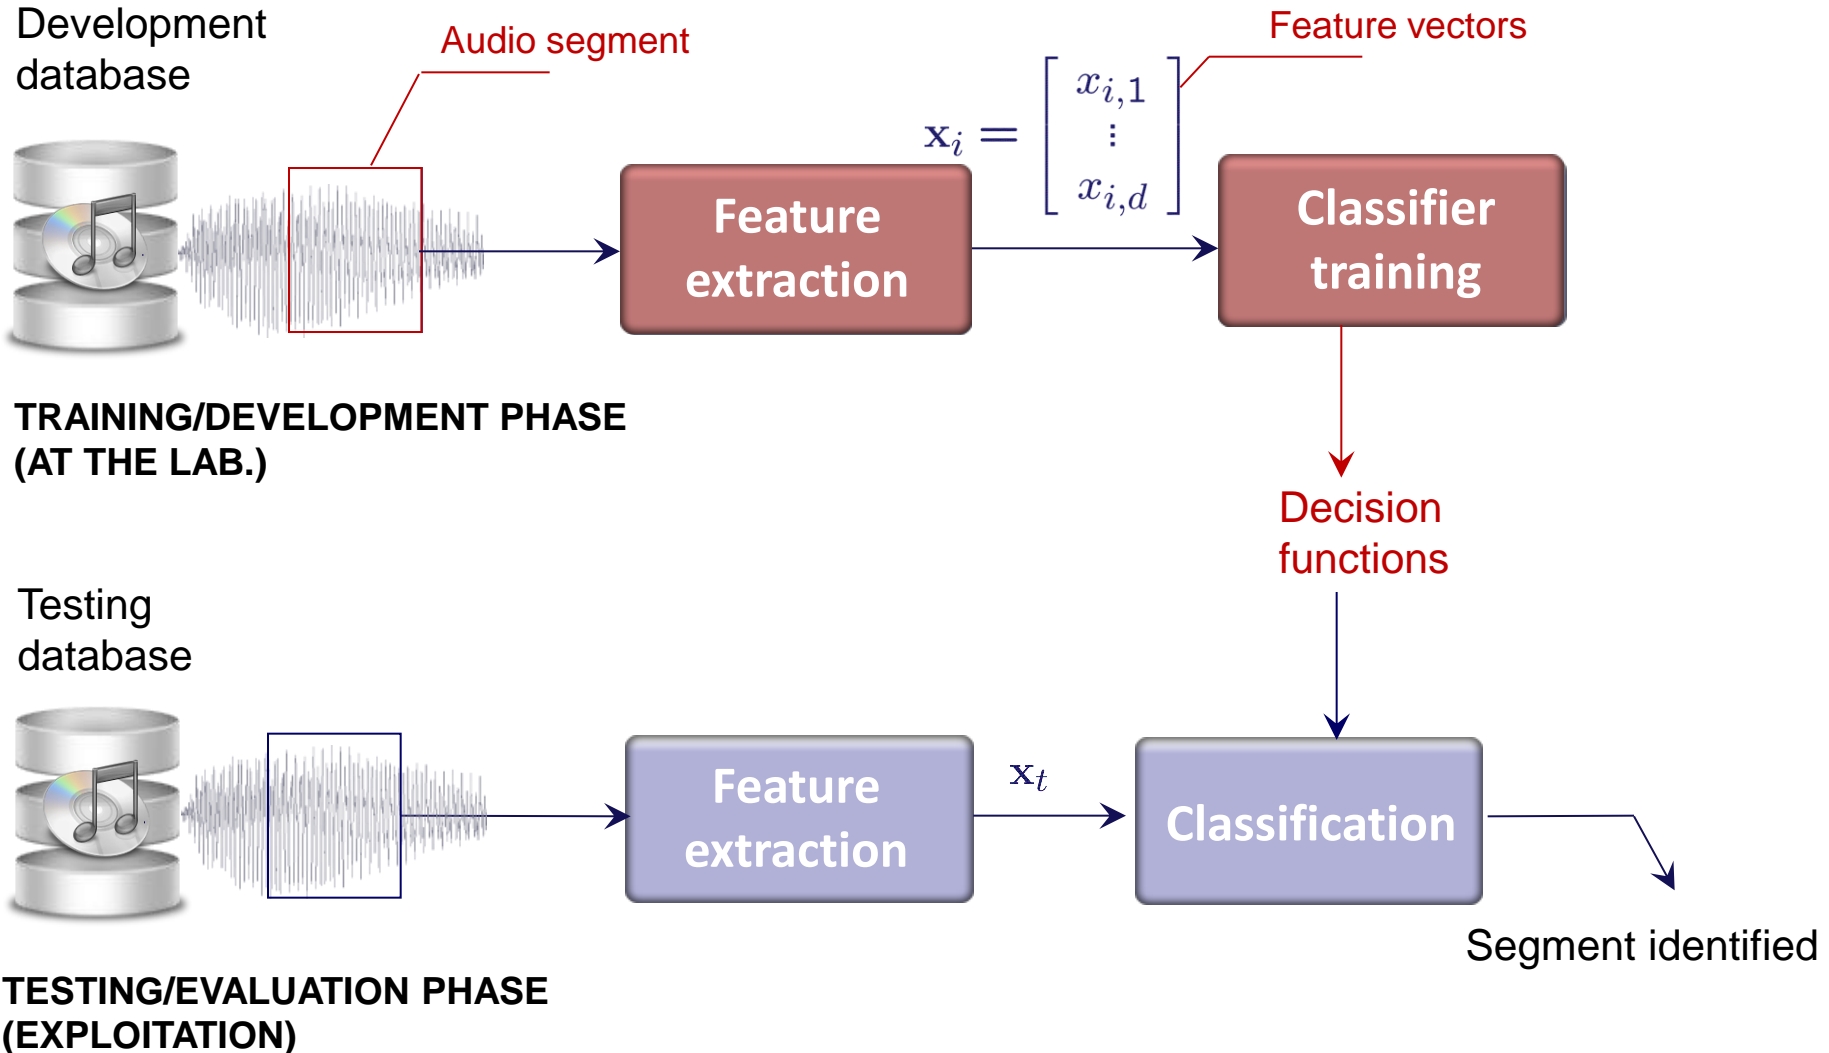


→ At the heart of all components: a **classification** task (supervised or unsupervised)

# General classification architecture




## ► Overview



# Classification architecture

## ► Feature extraction process



### Motivation:

- signal denoising/enhancement
- information rate reduction, eg. subsampling
- normalisation, eg.:

$$\tilde{s}(n) = s(n) - \bar{s}, \quad \bar{s} = \frac{1}{L} \sum_{n=0}^{L-1} s(n)$$

$$\hat{s}(n) = \frac{\tilde{s}(n)}{\max_n |\tilde{s}(n)|}$$

### Exercise

In Python:


- load an audio file;
- normalise it;
- visualise it.

Use  
*scipy.io.wavfile*





## ► Feature extraction process




→ Relies on audio **signal processing** techniques


# Audio signal analysis




## ► Short-Term analysis windows



*Drawing by J. Laroche, modified*



» **Static** temporal segmentation



» **Dynamic** temporal segmentation




- **Temporal features:** extracted directly from the waveform samples
- **Spectral features:** extracted from a frequential representation of the signal
- **Perceptual features:** extracted using a perceptual representation based on **psychoacoustic** considerations



## Zero Crossing Rates

$$\frac{1}{2} \sum_{n=2}^N |\text{sign}(x_n) - \text{sign}(x_{n-1})|$$

*Characterises noisy  
and transient sections*







## ► Discrete Fourier Transform

$$X_k = \sum_{n=0}^{N-1} x_n \exp(-j2\pi \frac{k}{N}n),$$

$$x_n = \frac{1}{N} \sum_{k=0}^{N-1} X_k \exp(j2\pi \frac{k}{N}n)$$

$$|X_k|$$



In practice: computed using the **Fast Fourier Transform (FFT)**

# Discrete Fourier Transform (DFT)



## ► Important properties

- Being a **discrete time** Fourier Transform, the DFT is **periodic**, with period 1 (in reduced frequency  $f = \frac{f}{f_s}$  ;  $f_s$  : sampling frequency)
- For signals  $x(n)$  and  $y(n)$  ;  $n \in \{0, \dots, N - 1\}$




Property	Numerical series	DFT
Linearity	$\{ax(n) + by(n)\}$	$\{aX(k) + bY(k)\}$
Hermitian symmetry	$x(n)$ <b>real</b>	$X(k) = X^*(-k)$
Time translation	$x(n - n_0)$	$X(k)e^{-\frac{2j\pi k}{N}n_0}$
Convolution (filtering)	$x(n) \star y(n)$ $\triangleq \sum_k x(k)y(n - k)$	$X(k)Y(k)$
Conjugation	$\{x^*(n)\}$	$\{X^*(-k)\}$



# Spectral analysis




► Violin excerpt: 20-ms overlapping windows (  $s_r = 44.1\text{kHz}$  ;  $N = 882$  samples )



# Spectral analysis

## ► Spectral analysis by Short-Term Fourier Transform (STFT)



*Drawing by J. Laroche*




# Spectral analysis



## ► Spectrogram

C note (262 Hz) produced by a **piano** and a **violin**



From M. Mueller & al. « *Signal Processing for Music Analysis*, *IEEE Trans. On Selected topics in Signal Processing* », October 2011.



## ► Exercise

In Python:

- Compute short-term spectra of an audio signal using FFT
- At home: compute and display spectrogram
- Use
  - » **scipy.io.wavfile**
  - » **pylab.specgram**
- Compare to hand crafted spectrogram obtained with:
  - » **scipy.fftpack** and **pylab.imshow**


# Spectral analysis



## ► Limitations of the spectrogram representation

- Large representation
  - » Typically 512 coefs every 10 ms
  - » High dimensionality
- Much detail
  - » Redundant representation
  - » High-level features (pitch, vibrato, timbre) are not highlighted

→ Still a low-level **representation**, not yet a **model**




# The source-filter model



- Distinction between:

- » **source**: excitation → fine spectral structure

- » **filter**: resonator → coarse structure





## ► Principle

- Source-filter model:  $y(n) = x(n) * h(n)$
- In the frequency domain:  $Y(f) = X(f)H(f)$   
 $\log |Y(f)| = \log |X(f)| + \log |H(f)|$
- By inverse DFT:  $c_y(q) = c_x(q) + c_h(q)$

where  $c_y(q) = \text{iDFT}[\log |Y(f)|]$ : real **cepstrum** definition


→ **deconvolution** is thus achieved: filter is separated from excitation

- First few cepstral coefficients
  - » low **quefrency**: “slow iDFT waves”
  - » represent the filter → **spectral envelope**
- Next coefficients represent the source → **fine spectral structure**





# Cepstral representations


## ► MFCC: Mel Frequency Cepstral Coefficients



Audio frame




Magnitude spectrum





Triangular filter banc  
in **Mel** scale

**Mel scale**




# Cepstral representations

## ► MFCC: Mel Frequency Cepstral Coefficients



# Cepstral representations

## ► MFCC: Mel Frequency Cepstral Coefficients

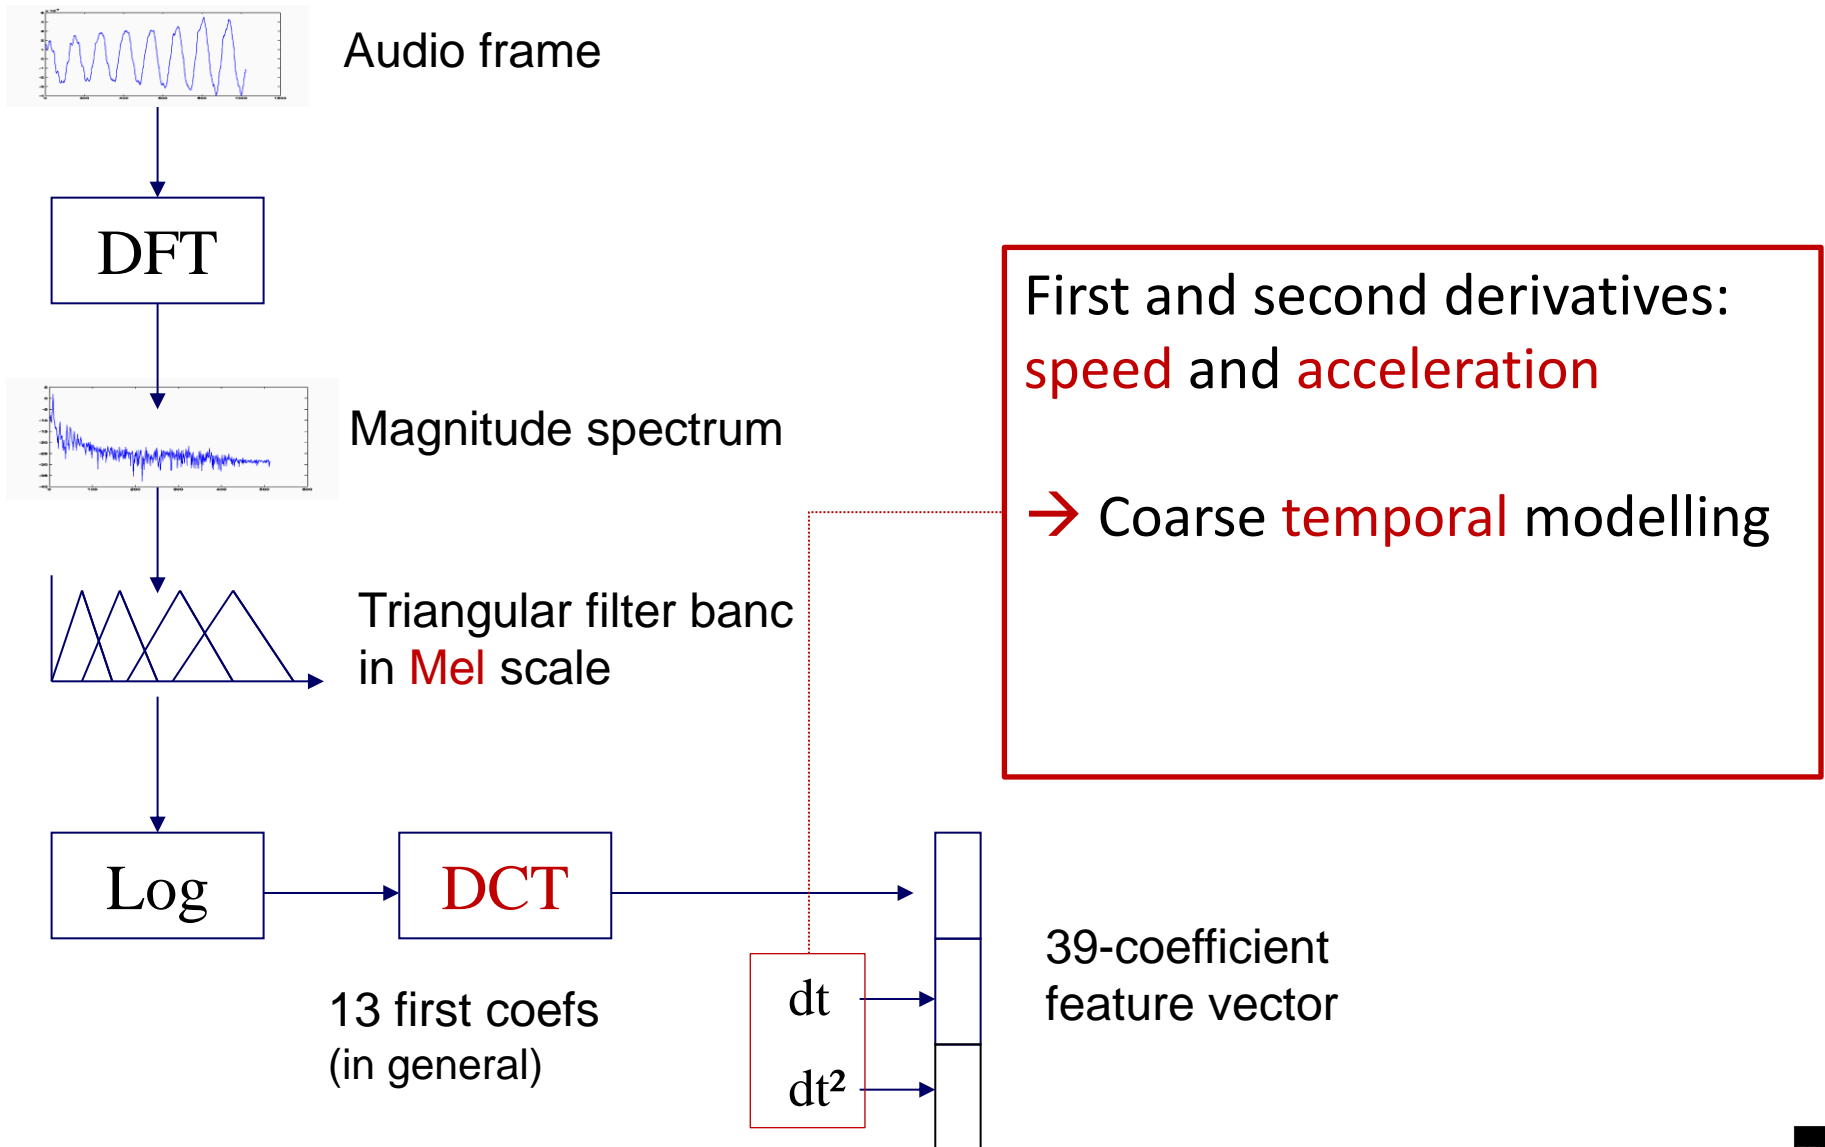


### Discrete Cosine Transform:

- nice **decorrelation** properties (like **PCA**)
- yields **diagonal** covariance matrices

# Cepstral representations

## ► MFCC: Mel Frequency Cepstral Coefficients





► ... very popular!

- In speech applications:
  - » Well justified: source-filter model makes sense
  - » Nice properties from a statistical modelling viewpoint: **decorrelation**
  - » Effective: state-of-the-art features for speaker and speech tasks
- In general audio classification:
  - » “Source-filter” model does not always hold
  - » Still, MFCCs work well in practice! they are the **default choice**



## ► Exercise

- Use librosa to extract MFCCs from an audio file
- Visualise the result

# Other spectral features: **spectral moments**



- Ordre 1: **Centre de Gravité Spectral** (centroïde spectral)

$$CGS = \frac{\sum_{k=1}^N k \cdot |X_k|}{\sum_{k=1}^N |X_k|}$$


- CGS élevé: son brillant
- CGS faible: son chaud, rond

- Ordre 2 : **Rayon de Giration Spectral**

$$RGS = \sqrt{\frac{\sum_{k=1}^N (k - CGS)^2 \cdot |X_k|}{\sum_{k=1}^N |X_k|}}$$

- RGS faible, le timbre est "compact"

- Ordres 3,4 également utilisés...







## Fréquence de coupure

fréquence  $F_c$  au dessous de laquelle  
85% de la distribution spectrale est  
concentrée


$$\sum_{k=1}^{F_c} |X_k| = 0.85 \times \sum_{k=1}^N |X_k|$$

## Platitude spectrale

mesurée par sous-bandes  $sb$  (**MPEG7 ASF**)

$$ASF(sb) = \frac{(\prod_{k \in sb} X_k)^{\frac{1}{K_{sb}}}}{\frac{1}{K_{sb}} \sum_{k \in sb} X_k}$$

Spectre plat :  $ASF \nearrow$ ,  $0 < ASF < 1$

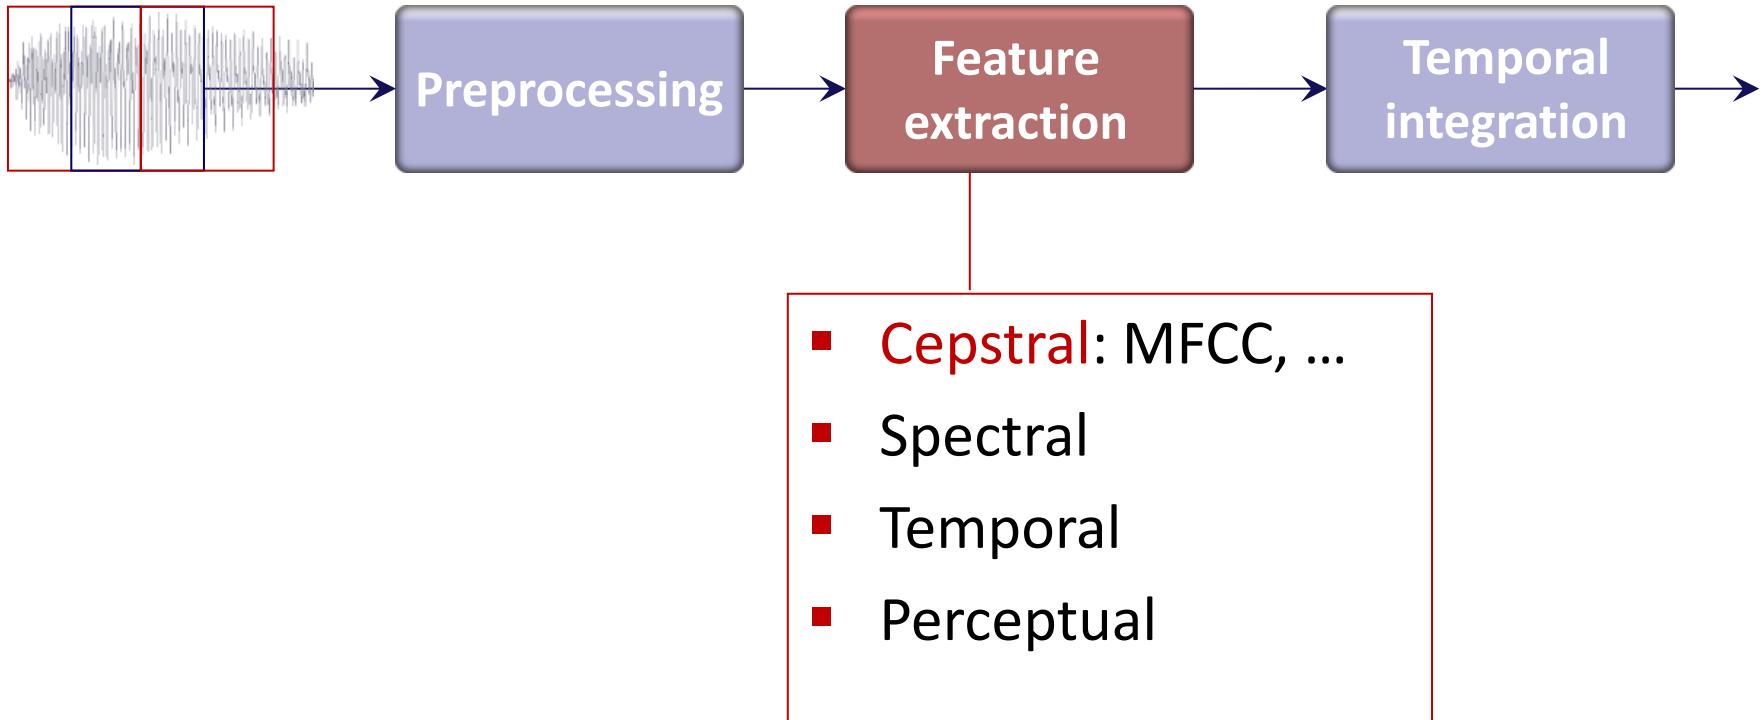


## Flux spectral (variation temporelle du contenu spectral)

$$Flux = \sum_{k=1}^N (|X_k(m)| - |X_k(m-1)|)^2$$




## ► Feature extraction process



## Which features to use?



## ► Feature extraction process




Which features to use for a given task?

- Use intuition/expert knowledge
- Use automatic **feature selection** algorithms
- Alternatively, use **feature learning**

# Classification architecture

## ► Feature extraction process




$$x_\tau = g\{x[k_\tau], \dots, x[k_\tau + L - 1]\}$$

$$\text{eg., } x_\tau = \text{mean}\{x[k_\tau], \dots, x[k_\tau + L - 1]\}$$



# Temporal integration

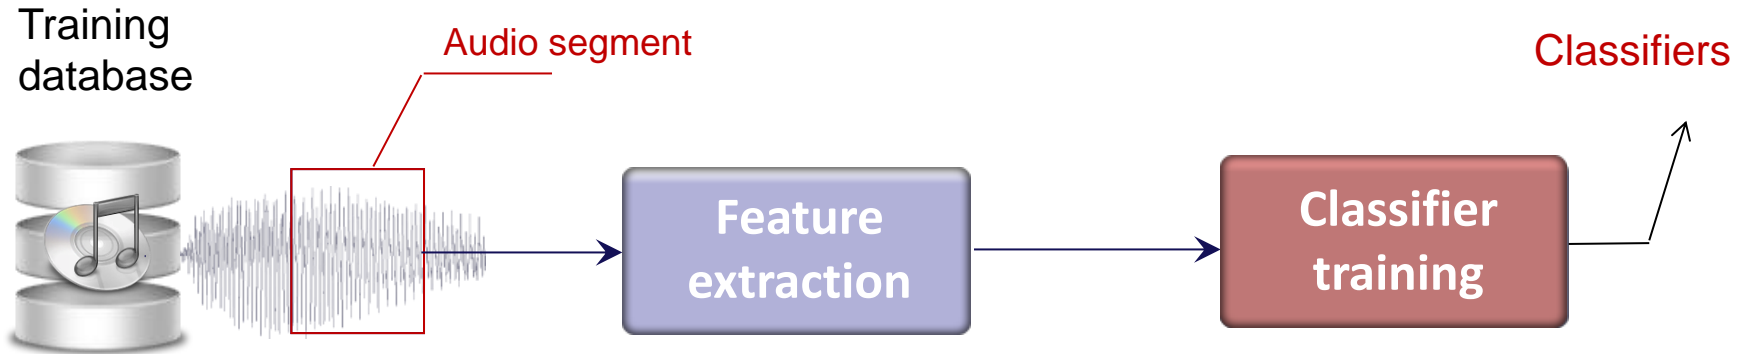
## ► At the feature level



- smoothing to improve robustness
- synchronising features extracted from different temporal horizons
- capturing temporal evolution of features

# Classification architecture

## ► Classifier training



**Training data:** assembled from all available audio instances

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_i^T \\ \vdots \\ \mathbf{x}_l^T \end{pmatrix} = \begin{pmatrix} x_{1,1} & \dots & x_{1,j} & \dots & x_{1,d} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ x_{i,1} & \dots & x_{i,j} & \dots & x_{i,d} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ x_{l,1} & \dots & x_{l,j} & \dots & x_{l,d} \end{pmatrix}, \quad \mathbf{y} = \begin{pmatrix} y_1 \\ \vdots \\ y_i \\ \vdots \\ y_l \end{pmatrix}$$

**Training examples**

**Class labels**

Unknown in non-supervised problems



## ■ **Books**

- » (Rabiner, 93) L. R. Rabiner, *Fundamentals of Speech Processing*. PTR Prentice-Hall, Inc., 1993.
- » (Ben Gold et al., 2011) B. Gold, N. Morgan, and D. Ellis, *Speech and Audio Signal Processing: Processing and Perception of Speech and Music*. John Wiley & Sons, 2011.
- » (Klapuri & Davy, 2010 ) A. Klapuri and M. Davy, *Signal Processing Methods for Music Transcription*, vol. 1. Springer, 2006.
- » (Kompatsiaris et al., 2012) *TV Content Analysis: Techniques and Applications (Multimedia Computing, Communication and Intelligence)*, Yiannis Kompatsiaris (Editor), Bernard Merialdo (Editor), Shiguo Lian (Editor). Taylor & Francis, 2012.
- » (Troncy et al. 2011) Raphael Troncy (Editor), Benoit Huet (Co-Editor), Simon Schenk (Co-Editor). *Multimedia Semantics: Metadata, Analysis and Interaction*. John Wiley & Sons, 2011.

## ■ **Articles and others**

- » (Peeters, 2003) G. Peeters, “A large set of audio features for sound description (similarity and classification) in the CUIDADO project,” IRCAM, 2004.
- » Software: librosa, YAAFE, MARSYAS, Sonic Annotator, MIR toolbox, .openSMILE, ...